




CPSC
4420/6420:
Artificial Intelligence

Submitted By:
Charanjit Singh
Parampreet Singh

Homework Project 3 (due Wednesday, November 09)

Player 1 Wins



.

In this programming assignment, you will design agents for the zero-sum game Connect Four, in which two players take turns dropping a colored disc from the top into a 7x6 grid. The game is won by forming a horizontal, vertical, or diagonal line of four same colored discs. See [here](#) for more details.

Basic Requirements

After downloading the attached source code, you should be able to display the game environment and select different game modes including Agent vs Agent, and Agent vs Human by typing the following at the command line:

```
python connect4.py
```

You can run the human-vs-human version to review how the game is played. Your task is to modify `connect4.py` in order to implement three of algorithms covered in class: minimax, alpha-beta pruning, and expectimax. Note that as different algorithms differ only in the dispatching mechanism, if you have minimax working correctly, implementing the rest of the methods should be straightforward.

Question 1 (4 points) Minimax

Modify the `minimax` function in `connect4.py` to implement depth-limited minimax search. As it's not feasible to search the entire game tree, your code should limit the search to an arbitrary depth by using the GUI. Score the leaves of your minimax tree with the

supplied `evaluate` function in order to treat them as terminal nodes.

Ans 1) [Click here for minimax code](#)

Question 2 (3 points) Alpha-Beta Pruning

Modify the `alphabeta` function in `connect4.py` to use alpha-beta pruning and allow more efficiently exploration of the minimax tree. You should be able to see a speed-up as the depth of the tree increases.

Ans 2) [Click here for minimax with alpha beta pruning code](#)

Question 3 (3 points) Expectimax

Minimax and alpha-beta assume that MAX plays against an adversary who makes optimal decisions. Modify the `expectimax` function in `connect4.py` to model probabilistic behavior of opponents that may make suboptimal decisions. To do so, you will replace MIN nodes (Agent2 or Human) with chance nodes. To simplify your code, assume you will only be running against an adversary which chooses actions uniformly at random.

Ans 3) [Click here for expectimax code](#)

Important notes:

1. Depending on your GUI selection, the game can be played between two agents, or an agent and a human player, or an agent and a random player. When an agent needs to make its next move, it runs an adversarial search as a MAX player and its opponent is considered to be a MIN (or a CHANCE) player depending on the search algorithm that you run. A random player simply makes a random valid move.
2. The minimax algorithm always considers that the adversary tries to minimize the score of the MAX player that initiated the game search. The adversary never considers its own score at all during this process. Therefore, when evaluating the utilities of the nodes at the maximum tree depth, the evaluation should always be made from MAX's point of view.
3. The pseudocode provided in the slides only returns the best utility value. However, here, you need to select the policy, i.e., the action that is associated with this value. To do so, you should consider all valid actions for the MAX player at the root of the tree, and return the one that leads to the best value.
4. We have provided you with a decent evaluation function that allows for effective depth limited search. Feel free to edit the `evaluate` function, if you are able to come up with a better one. A better function means lower computation time and/or higher winning percentage as compared to ours for a fixed depth. In case you're up to the task, please include a brief README explaining your strategy.

Submission

Please submit your modified `connect4.py` file to Project/Assignment 3 on [Canvas](#). You can work [in pairs](#) if you want to. If you do so, include both names in your report.

Getting Help

If you get stuck, please do not hesitate to contact the TA and the instructor for help, and stop by during office hours. We also encourage you to post questions and initiate discussions on [Canvas](#) and [discord](#). Your colleagues are also there to help you.

```

1 # connect4.py
2 # -----
3 # Licensing Information: You are free to use or
  extend these projects for
4 # educational purposes provided that (1) you do not
  distribute or publish
5 # solutions, (2) you retain this notice, and (3) you
  provide clear
6 # attribution to Clemson University and the authors.
7 #
8 # Authors: Pei Xu (peix@g.clemson.edu) and Ioannis
  Karamouzas (ioannis@g.clemson.edu)
9 #
10
11 """
12 In this assignment, the task is to implement the
  minimax algorithm with depth
13 limit for a Connect-4 game.
14
15 To complete the assignment, you must finish these
  functions:
16     minimax (line 196), alphabeta (line 237), and
  expectimax (line 280)
17 in this file.
18
19 In the Connect-4 game, two players place discs in a 6
  -by-7 board one by one.
20 The discs will fall straight down and occupy the
  lowest available space of
21 the chosen column. The player wins if four of his or
  her discs are connected
22 in a line horizontally, vertically or diagonally.
23 See https://en.wikipedia.org/wiki/Connect\_Four for
  more details about the game.
24
25 A Board() class is provided to simulate the game
  board.
26 It has the following properties:
27     b.rows           # number of rows of the game
  board
28     b.cols           # number of columns of the game

```

```

28 board
29     b.PLAYER1      # an integer flag to represent
    the player 1
30     b.PLAYER2      # an integer flag to represent
    the player 2
31     b.EMPTY_SLOT   # an integer flag to represent an
    empty slot in the board;
32
33 and the following methods:
34     b.terminal()    # check if the game is
    terminal
35                     # terminal means draw or
    someone wins
36
37     b.has_draw()    # check if the game is a
    draw
38
39     w = b.who_wins() # return the winner of
    the game or None if there
40                     # is no winner yet
41                     # w should be in [b.
    PLAYER1,b.PLAYER2, None]
42
43     b.occupied(row, col) # check if the slot at
    the specific location is
44                     # occupied
45
46     x = b.get(row, col) # get the player
    occupying the given slot
47                     # x should be in [b.
    PLAYER1, b.PLAYER2, b.EMPTY_SLOT]
48
49     row = b.row(r)    # get the specific row of
    the game described using
50                     # b.PLAYER1, b.PLAYER2
    and b.EMPTY_SLOT
51
52     col = b.column(r) # get a specific column
    of the game board
53
54     b.placeable(col)  # check if a disc can be

```

```

54 placed at the specific
55                                     # column
56
57     b.place(player, col)    # place a disc at the
    specific column for player
58     # raise ValueError if the specific column
    does not have available space
59
60     new_board = b.clone()    # return a new board
    instance having the same
61                                     # disc placement with b
62
63     str = b.dump()           # a string to describe
    the game board using
64                                     # b.PLAYER1, b.PLAYER2
    and b.EMPTY_SLOT
65 Hints:
66     1. Depth-limited Search
67         We use depth-limited search in the current
    code. That is we
68     stop the search forcefully, and perform
    evaluation directly not only when a
69     terminal state is reached but also when the
    search reaches the specified
70     depth.
71     2. Game State
72         Three elements decide the game state. The
    current board state, the
73     player that needs to take an action (place a disc
    ), and the current search
74     depth (remaining depth).
75     3. Evaluation Target
76         The minimax algorithm always considers that
    the adversary tries to
77     minimize the score of the max player, for whom
    the algorithm is called
78     initially. The adversary never considers its own
    score at all during this
79     process. Therefore, when evaluating nodes, the
    target should always be
80     the max player.

```

```

81     4. Search Result
82         The pseudo code provided in the slides only
           returns the best utility value.
83         However, in practice, we need to select the
           action that is associated with this
84         value. Here, such action is specified as the
           column in which a disc should be
85         placed for the max player. Therefore, for each
           search algorithm, you should
86         consider all valid actions for the max player,
           and return the one that leads
87         to the best value.
88
89     """
90
91     # use math library if needed
92     import math
93
94     def get_child_boards(player, board):
95         """
96         Generate a list of successor boards obtained by
           placing a disc
97         at the given board for a given player
98
99         Parameters
100        -----
101        player: board.PLAYER1 or board.PLAYER2
102               the player that will place a disc on the
           board
103        board: the current board instance
104
105        Returns
106        -----
107        a list of (col, new_board) tuples,
108        where col is the column in which a new disc is
           placed (left column has a 0 index),
109        and new_board is the resulting board instance
110        """
111        res = []
112        for c in range(board.cols):
113            if board.placeable(c):

```

```

114         tmp_board = board.clone()
115         tmp_board.place(player, c)
116         res.append((c, tmp_board))
117     return res
118
119
120 def evaluate(player, board):
121     """
122     This is a function to evaluate the advantage of
123     the specific player at the
124     given game board.
125
126     Parameters
127     -----
128     player: board.PLAYER1 or board.PLAYER2
129             the specific player
130     board: the board instance
131
132     Returns
133     -----
134     score: float
135            a scalar to evaluate the advantage of the
136            specific player at the given
137            game board
138     """
139     adversary = board.PLAYER2 if player == board.
140     PLAYER1 else board.PLAYER1
141     # Initialize the value of scores
142     # [s0, s1, s2, s3, --s4--]
143     # s0 for the case where all slots are empty in a
144     4-slot segment
145     # s1 for the case where the player occupies one
146     slot in a 4-slot line, the rest are empty
147     # s2 for two slots occupied
148     # s3 for three
149     # s4 for four
150     score = [0]*5
151     adv_score = [0]*5
152
153     # Initialize the weights
154     # [w0, w1, w2, w3, --w4--]

```

```

150     # w0 for s0, w1 for s1, w2 for s2, w3 for s3
151     # w4 for s4
152     weights = [0, 1, 4, 16, 1000]
153
154     # Obtain all 4-slot segments on the board
155     seg = []
156     invalid_slot = -1
157     left_revolved = [
158         [invalid_slot]*r + board.row(r) + \
159         [invalid_slot]*(board.rows-1-r) for r in
range(board.rows)
160     ]
161     right_revolved = [
162         [invalid_slot]*(board.rows-1-r) + board.row(
r) + \
163         [invalid_slot]*r for r in range(board.rows)
164     ]
165     for r in range(board.rows):
166         # row
167         row = board.row(r)
168         for c in range(board.cols-3):
169             seg.append(row[c:c+4])
170     for c in range(board.cols):
171         # col
172         col = board.col(c)
173         for r in range(board.rows-3):
174             seg.append(col[r:r+4])
175     for c in zip(*left_revolved):
176         # slash
177         for r in range(board.rows-3):
178             seg.append(c[r:r+4])
179     for c in zip(*right_revolved):
180         # backslash
181         for r in range(board.rows-3):
182             seg.append(c[r:r+4])
183     # compute score
184     for s in seg:
185         if invalid_slot in s:
186             continue
187         if adversary not in s:
188             score[s.count(player)] += 1

```



```

189         if player not in s:
190             adv_score[s.count(adversary)] += 1
191         reward = sum([s*w for s, w in zip(score, weights
192         )])
192         penalty = sum([s*w for s, w in zip(adv_score,
193         weights)])
193         return reward - penalty
194
195     #Depth limited minimax search
196
197 def minimax(player, board, depth_limit):
198
199     """
200     Minimax algorithm with limited search depth.
201
202     Parameters
203     -----
204     player: board.PLAYER1 or board.PLAYER2
205           the player that needs to take an action (
206           place a disc in the game)
207     board: the current game board instance
208     depth_limit: int
209           the tree depth that the search algorithm
210           needs to go further before stopping
211     max_player: boolean
212
213     Returns
214     -----
215     placement: int or None
216           the column in which a disc should be placed
217           for the specific player
218           (counted from the most left as 0)
219           None to give up the game
220     """
221     max_player = player
222     next_player = board.PLAYER2 if player == board.
223     PLAYER1 else board.PLAYER1
224     placement = None
225
226     ### Please finish the code below
227     #####

```

```

223 #####
224 #####
225     def value(player, board, depth_limit):
226         if board.terminal() or depth_limit == 0:
227             return (evaluate(player, board), None)
228         elif player == max_player:
229             return max_value(player, board,
230 depth_limit)
231         elif player == next_player:
232             return min_value(player, board,
233 depth_limit)
234
235     def max_value(player, board, depth_limit):
236         l = []
237         for c, b in get_child_boards(player, board):
238             v, a = value(next_player, b, depth_limit
239 -1)
240             l.append((v, c))
241         return max(l)
242
243     def min_value(player, board, depth_limit):
244         l = []
245         for c, b in get_child_boards(player, board):
246             v, a = value(max_player, b, depth_limit-
247 1)
248             l.append((v, c))
249         return min(l)
250
251     score, placement = value(max_player, board,
252 depth_limit)
253
254 #####
255 #####
256     return placement
257
258 #Alpha beta pruning of minimax search for efficiency
259 def alphabeta(player, board, depth_limit):
260     """
261     Minimax algorithm with alpha-beta pruning.
262
263     Parameters

```

```

257      -----
258      player: board.PLAYER1 or board.PLAYER2
259      the player that needs to take an action (
place a disc in the game)
260      board: the current game board instance
261      depth_limit: int
262      the tree depth that the search algorithm
needs to go further before stopping
263      alpha: float
264      beta: float
265      max_player: boolean
266
267
268      Returns
269      -----
270      placement: int or None
271      the column in which a disc should be placed
for the specific player
272      (counted from the most left as 0)
273      None to give up the game
274      """
275      max_player = player
276      next_player = board.PLAYER2 if player == board.
PLAYER1 else board.PLAYER1
277      alpha = -math.inf
278      beta = math.inf
279      placement = None
280
281      ### Please finish the code below
#####
282 #####
#####
283      def value(player, board, depth_limit, alpha,
beta):
284          if board.terminal():
285              return (evaluate(player, board), None)
286          elif depth_limit == 0:
287              return (evaluate(player, board), None)
288          elif player == max_player:
289              return max_value(player, board,
depth_limit, alpha, beta)

```

```

290         elif player == next_player:
291             return min_value(player, board,
292                               depth_limit, alpha, beta)
293     def max_value(player, board, depth_limit, alpha
294                   , beta):
295         l = []
296         for c, b in get_child_boards(player, board):
297             v, a = value(next_player, b, depth_limit
298                           -1, alpha, beta)
299             if v >= beta:
300                 return (v, c)
301             l.append((v, c))
302             alpha = max(alpha, v)
303         return max(l)
304     def min_value(player, board, depth_limit, alpha
305                   , beta):
306         l = []
307         for c, b in get_child_boards(player, board):
308             v, a = value(max_player, b, depth_limit-
309                           1, alpha, beta)
310             if v <= alpha:
311                 return (v, c)
312             l.append((v, c))
313             beta = min(beta, v)
314         return min(l)
315     score, placement = value(max_player, board,
316                               depth_limit, -math.inf, math.inf)
317     #####
318     #####
319     return placement
320
321     #Expectimax search with uniform chance of selecting
322     available actions
323
324     def expectimax(player, board, depth_limit):
325         """
326         Expectimax algorithm.
327         We assume that the adversary of the initial

```

```

322 player chooses actions
323     uniformly at random.
324     Say that it is the turn for Player 1 when the
function is called initially,
325     then, during search, Player 2 is assumed to pick
actions uniformly at
326     random.
327
328     Parameters
329     -----
330     player: board.PLAYER1 or board.PLAYER2
331     the player that needs to take an action (
place a disc in the game)
332     board: the current game board instance
333     depth_limit: int
334     the tree depth that the search algorithm
needs to go before stopping
335     max_player: boolean
336
337     Returns
338     -----
339     placement: int or None
340     the column in which a disc should be placed
for the specific player
341     (counted from the most left as 0)
342     None to give up the game
343     """
344     max_player = player
345     next_player = board.PLAYER2 if player == board.
PLAYER1 else board.PLAYER1
346     placement = None
347
348     ### Please finish the code below
#####
349
#####
#####
350     def value(player, board, depth_limit):
351         if board.terminal() or depth_limit == 0:
352             return (evaluate(player, board), None)
353         elif player == max_player:

```

```

354         return max_value(player, board,
    depth_limit)
355     elif player == next_player:
356         return exp_value(player, board,
    depth_limit)
357
358     def prob_action(player, board):
359         possible_Actions = 0
360         for c in range(board.cols):
361             if board.placeable(c):
362                 possible_Actions += 1
363         return (1/possible_Actions)
364
365     def max_value(player, board, depth_limit):
366         l = []
367         for c, b in get_child_boards(player, board):
368             v, a = value(next_player, b, depth_limit
    - 1)
369             l.append((v, c))
370         return max(l)
371
372     def exp_value(player, board, depth_limit):
373         l = 0
374         for c, b in get_child_boards(player, board):
375             v, a = value(max_player, b, depth_limit
    - 1)
376             l = prob_action(player, board)*v + l
377         return (l, None)
378
379     score, placement = value(max_player, board,
    depth_limit)
380
381     #####
    #####
382     return placement
383
384 if __name__ == "__main__":
385     from utils.app import App
386     import tkinter
387

```

```
388     algs = {
389         "Minimax": minimax,
390         "Alpha-beta pruning": alphabeta,
391         "Expectimax": expectimax
392     }
393
394     root = tkinter.Tk()
395     App(algs, root)
396     root.mainloop()
397
```